

```
In [1]: # import libraries
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

## Gradient Descent (in general)


It is an optimization algorithm to find the minimum of a function. We start with a random point on the function and move in the negative direction of the gradient of the function to reach the local/global minima.

Lets say the function is  $y=(x-5)^2$

```
In [2]: X = np.arange(-5,20)
y = (X-5)**2
print(X)
print(y)
```

```
[-5 -4 -3 -2 -1  0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18
 19]
[100  81  64  49  36  25  16   9   4   1   0   1   4   9  16  25  36  49
  64  81 100 121 144 169 196]
```

```
In [3]: # Lets visualize the function
plt.plot(X,y)
plt.ylabel('f(x)=(x+5)²')
plt.xlabel('X')
plt.show()
```



## Gradient Descent formula

$$\theta^t = \theta^{t-1} - \eta \frac{\partial J}{\partial \theta^{t-1}}$$

where

$\theta^t$  -> next step or next point

$\theta^{t-1}$  -> previous step or previous point

$\eta$  -> Learning rate

J -> function which we have to minimise

$\frac{\partial J}{\partial \theta^{t-1}}$  -> partial derivative of that function or called Gradient

```
In [4]: # Lets find the minimum of that function y = (X-5)**2
# Lets start from x = 15

X = 15 # starting point
lr = 0.1 # learning rate

for i in range(100): # Lets take 100 iteration
    grad = 2*(X-5) # derivative of that function or gradient
    X=X-lr*grad # applying formula
print(X)
```

5.000000002037036

## conclusion


- So approximately 5 is our global minima for the function  $y = (X-5)^2$ .
- That means at the value of  $X = 5$  our function is minimum.

In [5]: *# Let's visualize how we iteratively converge to the local minima.*

```
X = 15
lr = 0.1

for i in range(100):
    grad = 2*(X-5)
    X=X-lr*grad
    y = (X-5)**2
    plt.scatter(X,y)
print(X)
```

5.000000002037036



## Gradient Descent in Linear Regression

In [6]: *# Lets take a dataset called housing.csv*

```
housing = pd.read_csv(r'E:\linkedin post project\Gradient Descent\housing.csv')
housing.head()
```

Out[6]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	13300000	7420	4	2	3	1	0	0	0	1	2	1	furnished
1	12250000	8860	4	4	4	1	0	0	0	1	3	0	furnished
2	12250000	9960	3	2	2	1	0	1	0	0	2	1	semi-furnished
3	12215000	7500	4	2	2	1	0	1	0	1	3	1	furnished
4	11410000	7420	4	1	2	1	1	1	0	1	2	0	furnished

In [7]: *#Converting furnishingstatus column to binary column using get\_dummies*

```
status = pd.get_dummies(housing['furnishingstatus'],drop_first=True)
housing = pd.concat([housing,status],axis=1)
housing.drop(['furnishingstatus'],axis=1,inplace=True)
```

```
In [8]: # Normalising the data
housing = (housing - housing.mean())/housing.std()
housing.head()
```

```
Out[8]:
```


	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	semi-furnished	unfurnished
0	4.562174	1.045766	1.402131	1.420507	1.376952	0.405251	-0.464888	-0.733865	-0.219063	1.471267	1.518299	1.803284	-0.844113	-1
1	4.000809	1.755397	1.402131	5.400847	2.529700	0.405251	-0.464888	-0.733865	-0.219063	1.471267	2.876950	-0.553526	-0.844113	-1
2	4.000809	2.216198	0.047235	1.420507	0.224204	0.405251	-0.464888	1.360148	-0.219063	-0.678439	1.518299	1.803284	1.182502	-1
3	3.982096	1.082630	1.402131	1.420507	0.224204	0.405251	-0.464888	1.360148	-0.219063	1.471267	2.876950	1.803284	-0.844113	-1
4	3.551716	1.045766	1.402131	-0.560663	0.224204	0.405251	2.147110	1.360148	-0.219063	1.471267	1.518299	-0.553526	-0.844113	-1

## Simple Linear Regression

```
In [9]: # Lets say we want to predict the house price using the area attributes
# Assign feature variable or independent variable X
X = housing['area']

# Assign response variable or dependent variable to y
y = housing['price']
```

```
In [10]: # Lets visualize relationship between the features and the response using scatterplots
sns.scatterplot(X,y)
plt.show()
```



For linear regression we use a cost function known as the mean squared error or MSE.

Formula

$$MSE = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - \hat{y})^2$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - h_{\theta}x^{(i)})^2$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - (mx^{(i)} + c))^2$$

where

$n$  = number of items

$y^{(i)}$  = actual point

$\hat{y} = h_{\theta}x^{(i)} = mx^{(i)} + c$  = predicted point

$m$  = slope

$c$  = intercept

So our goal is to find the best  $m$  and  $c$  value where our MSE should be minimum .

As per Gradient Descent Formula we have to differentiate our cost function first and then put that in to the formula .

**Partial Darivation of MSE w.r.t slope( $m$ ) and intercept( $c$ )**

$$\frac{\partial MSE}{\partial m} = -\frac{2}{n} \sum_{i=1}^n (y^{(i)} - mx^{(i)} - c)(x^{(i)})$$

$$\frac{\partial MSE}{\partial c} = -\frac{2}{n} \sum_{i=1}^n (y^{(i)} - mx^{(i)} - c)$$

```
In [11]: # now for applying gradient Descent we need our X,y variables as numpy array
X = np.array(X)
y = np.array(y)
```

```

In [12]: # Lets implement the gradient descent function
# Lets initialised current m and c to 0

def gradient_Descent_simple(X,y):
    m = 0
    c = 0
    n = float(len(y))
    iters = 1000 # take 1000 iteration
    learning_rate = 0.01
    df = pd.DataFrame(columns = ['m','c','cost'])# make a dataframe to keep track how the costs are minimising in each iteration
    for i in range(iters):
        y_pred = m*X+c # prediction values
        cost = sum([i**2 for i in (y-y_pred)])/n
        d_m = (-2/n)*sum(X*(y-y_pred))#darivative w.r.t m
        d_c = (-2/n)*sum(y-y_pred)#darivative w.r.t c
        # we are doing the derivation for minimising the cost
        # after darivation we have to update the m and c
        m = m - (learning_rate*d_m)
        c = c - (learning_rate*d_c)
        df.loc[i] = [m,c,cost] # keep track of each cost in each iteration
    df.reset_index().plot.line(x='index', y=['cost'])
    return f'final slope and intercept after 1000 iteration is {round(m,3),round(c,3)}'

```

```


In [13]: gradient_Descent_simple(X,y)

```

```


Out[13]: 'final slope and intercept after 1000 iteration is (0.536, 0.0)'

```



- We can clearly see that after 200 iterations the cost was not decreasing much which means we get our global minima.

```
In [14]: # Lets visualize how the line was fitted
y_pred = 0.536*X+0.0
plt.scatter(X,y)
plt.plot(X,y_pred,color = 'red')
plt.xlabel('area')
plt.ylabel('price')
plt.show()
```



#### Note

- learning rate is the speed at which we want to move towards negative of the gradient .
- it's always a good practice to choose a small value of learning rate and slowly move towards the negative of the gradient .

## Multiple Linear Regression

### Applying Gradient Descent for Multiple (>1) Features

```
In [15]: # Assigning feature variable X
# Lets now take 2 features area and bedrooms
X = housing[['area','bedrooms']]

# Assigning response variable y
y = housing['price']
```

#### Note

- when we have more than one features the model now fit a hyperplane instead of line .

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 \dots + \theta_n x_n$$

- The cost function here is slightly different from MSE we just dividing the MSE by half (1/2) to get a nice interpretation .

$$J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)}))^2$$

$$J(\theta_0, \theta_1, \dots, \theta_i) = \frac{1}{2n} \sum_{i=1}^n (y^{(i)} - (\theta_0 + \theta_1 x_1 + \theta_2 x_2 \dots + \theta_i x_i))^2$$

we have to first differentiate w.r.t. all thetas and then update each thetas .

$$\frac{\partial J(\theta_0, \theta_1, \dots, \theta_i)}{\partial \theta_0} = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) (x_0^{(i)})$$

$$\frac{\partial J(\theta_0, \theta_1, \dots, \theta_i)}{\partial \theta_1} = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) (x_1^{(i)})$$

...

$$\frac{\partial J(\theta_0, \theta_1, \dots, \theta_i)}{\partial \theta_i} = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) (x_i^{(i)})$$

- We need to minimise the cost function  $J(\theta)$  One way to do this is to use the batch gradient decent algorithm. In batch gradient decent, the values are updated in each iteration:

$$\theta_0 = \theta_0 - \alpha \frac{1}{n} \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) (x_0^{(i)})$$

$$\theta_1 = \theta_1 - \alpha \frac{1}{n} \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) (x_1^{(i)})$$

$$\theta_2 = \theta_2 - \alpha \frac{1}{n} \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) (x_2^{(i)})$$

...

$$\theta_n = \theta_n - \alpha \frac{1}{n} \sum_{i=1}^n (y^{(i)} - h_{\theta}(x^{(i)})) (x_n^{(i)})$$

- Now here we are taking the help of matrix multiplication because we are deal with more than one features .



```
In [16]: # Add a columns of 1s as an intercept to X.
# The intercept column is needed for convenient matrix representation of cost function

X['intercept'] = 1
X = X.reindex(['intercept', 'area', 'bedrooms'], axis=1)
X.head()
```

```
Out[16]:
```

	intercept	area	bedrooms
0	1	1.045788	1.402131
1	1	1.755397	1.402131
2	1	2.216196	0.047235
3	1	1.082630	1.402131
4	1	1.045788	1.402131

```
In [17]: # Convert X and y to arrays
import numpy as np
X = np.array(X)
y = np.array(y)
```

```
In [18]: # Theta is the vector representing coefficients (intercept, area, bedrooms)
theta = np.matrix(np.array([0,0,0]))
alpha = 0.01
iterations = 1000
```

```
In [19]: # define cost function
# takes in theta (current values of coefficients ), X and y
# returns total cost at current  $\theta_1, \theta_2, \theta_3$ 

def compute_cost(X, y, theta):
    return np.sum(np.square(np.matmul(X, theta) - y)) / (2 * len(y))
```


```
In [20]: # gradient descent
# takes in current X, y, Learning rate alpha, num_iters
# returns cost (notice it uses the cost function defined above)

def gradient_descent_multi(X, y, theta, alpha, iterations):
    theta = np.zeros(X.shape[1])
    n = len(X)
    df = pd.DataFrame( columns = ['coefficients', 'cost'])

    for i in range(iterations):
        cost = compute_cost(X, y, theta)
        derivative = (1/n) * np.matmul(X.T, np.matmul(X, theta) - y) # we are doing derivative for minimizing the cost
        theta = theta - alpha * derivative # here we have to update our theta means we are going to next step
        df.loc[i] = [theta, cost]
    df.reset_index().plot.line(x='index', y=['cost'])
    return df.tail(), f'final coefficients are {df.iloc[999,0]} '
```

```
In [21]: # print costs with various values of coefficients b0, b1, b2
gradient_descent_multi(X, y, theta, alpha, iterations)
```

```
Out[21]: (
              coefficients      cost
995 [3.3328691488600403e-16, 0.4916558741575689, 0... 0.314176
996 [3.3359248085608434e-16, 0.49165609672441263, ... 0.314176
997 [3.337635977993293e-16, 0.4916563172711537, 0... 0.314176
998 [3.3429324548008186e-16, 0.4916565358164564, 0... 0.314176
999 [3.3447658506285006e-16, 0.49165675237880896, ... 0.314176,
'final coefficients are [3.34476585e-16 4.91656752e-01 2.91844700e-01] ')
```



- We can clearly see that after 200 iterations the cost was not decreasing much which means we get our global minima.

## EXAMPLE 2

```
In [22]: # lets take a dataset called advertising.csv
advertising = pd.read_csv(r'E:\\advertising (1).csv')
advertising.head()
```

```
Out[22]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

```
In [23]: # Normalising the data
advertising = (advertising - advertising.mean())/advertising.std()
advertising.head()
```

```
Out[23]:
```


	TV	Radio	Newspaper	Sales
0	0.907425	0.979086	1.774493	1.548188
1	-1.194379	1.080097	0.867903	-0.894304
2	-1.512380	1.524037	1.779084	-0.905135
3	0.051919	1.214806	1.283185	0.858177
4	0.383198	-0.839507	1.278593	-0.215143

```
In [24]: # Assigning feature variable X
# Lets see how attribute 'TV' affects 'Sales'
X = advertising['TV']

# Assigning response variable y
y = advertising['Sales']


plt.scatter(X,y)
```

Out[24]: <matplotlib.collections.PathCollection at 0x2715a6628e0>




```
In [25]: # first convert X,y to numpy array
X = np.array(X)
y = np.array(y)
# call the function
gradient_Descent_simple(X,y)
```

Out[25]: 'final slope and intercept after 1000 iteration is (0.782, -0.0)'



In [26]: # Lets visualize how the Line was fitted

```
y_pred = 0.782*X-0.8
plt.scatter(X,y)
plt.plot(X,y_pred,color = 'red')
plt.xlabel('area')
plt.ylabel('price')
plt.show()
```



In [27]: # Assigning feature variable X  
# Lets now take all features  
X = advertising[['TV','Radio','Newspaper']]

```
# Assigning response variable y  
y = advertising['Sales']
```

In [28]: # Add a column of 1s as an intercept to X.  
# The intercept column is needed for convenient matrix representation of cost function

```
X['intercept'] = 1  
X = X.reindex(['TV','Radio','Newspaper'], axis=1)  
X.head()
```

Out[28]:

	TV	Radio	Newspaper
0	0.987425	0.979088	1.774493
1	-1.194379	1.080097	0.887903
2	-1.512380	1.524837	1.779084
3	0.051919	1.214808	1.283185
4	0.393198	-0.839507	1.278593

In [29]: # Convert X and y to arrays

```
import numpy as np  
X = np.array(X)  
y = np.array(y)
```


In [30]: # Theta is the vector representing coefficients (intercept, area, bedrooms)

```
theta = np.matrix(np.array([0,0,0,0]))  
alpha = 0.01  
iterations = 1000
```

```
In [31]: # call the function
gradient_descent_multi(X, y, theta, alpha, iterations)
```

```
Out[31]: (
              coefficients      cost
995 [0.7530275184009089, 0.5360440327094922, -0.00... 0.051138
996 [0.7530278865846947, 0.5360468290216902, -0.00... 0.051138
997 [0.7530282512106896, 0.5360496075026058, -0.00... 0.051138
998 [0.7530286123134639, 0.5360523682656025, -0.00... 0.051138
999 [0.75302896992725, 0.5360551114233569, -0.0038... 0.051138,
'final coefficients are [ 0.75302897 0.53605511 -0.0038943 ] ')

```



```
In [32]: print('@'*50+'THE END'+ '@'*50)
```

#####THE END#####